

# What is Machine Learning?

Oxford Spring School in Advanced Research Methods 2024

Dr Thomas Robinson, LSE

Day 1/5 (2024)

# Introduction

# This course

5 × 3 hour sessions, covering:

1. Introduction to ML (maximum likelihood estimation)
2. Regularised methods (bias variance trade-off)
3. Tree-based methods (hyperparameter tuning)
4. Neural networks (feature engineering)
5. Ensemble methods

The logic:

- ▶ Understand the underlying mechanics of parameter estimation
- ▶ Start from a modeling strategy we are familiar with...
- ▶ ... and move on to algorithmically more complicated cases
- ▶ Build on the same foundational concepts across each day

# What won't we cover?

- ▶ ML is a broad and contested domain
- ▶ We will not cover some topics:
  - ▶ Unsupervised methods
  - ▶ Clustering algorithms
  - ▶ (Text-specific ML models)

This course prioritises:

- ▶ Intuitive understanding of popular ML methods
- ▶ Transferability of fundamentals
- ▶ Relevance to “downstream” social science research problems

# Session structure

Each day will be a mixture of:

- ▶ Lecture content (approx. 2 hours)
  - ▶ Plenty of opportunities for questions
  - ▶ Time for us to think through some applied problems
- ▶ Coding walkthroughs (approx. 1 hour)
  - ▶ Conducted in R using RStudio
  - ▶ Hands on experience using different algorithms
- ▶ Self-guided learning
  - ▶ Applied readings using the methods we discuss in class
  - ▶ Completing and consolidating coding exercises

**Slides and code are available here:**

**[www.github.com/tsrobinson/oxss24](https://www.github.com/tsrobinson/oxss24)**

# Today's session

Goals:

1. Introduce the topic of machine learning
  - ▶ What is ML?
  - ▶ How do we distinguish ML from “traditional” statistics?
  - ▶ What sort of problems might we apply it to?
2. Introduce key conceptual distinctions we will make throughout the course
3. Introduce maximum likelihood estimation
  - ▶ A key way in which ML parameters are estimated
  - ▶ Build our own logistic regression estimator

What is machine learning?

## (Machine) learning and statistics

*“There are two cultures in the use of statistical modeling to reach **conclusions from data**. One assumes that the data are generated by a given stochastic data model. The other uses algorithmic models and treats the data mechanism as unknown” – Breiman 2001*

*“Statistical learning refers to a set of tools for modeling and understanding **complex** datasets. It is a recently developed area in statistics and blends with parallel developments in **computer science** and, in particular, machine learning” – James et al. 2013*

*“Machine learning is a subfield of **computer science** that is concerned with building algorithms which, to be useful, rely on a collection of examples of some phenomenon. . . ”  
– Burkov 2019*



# Machine learning

Expectation: I need a \$1m super computer

Reality: It runs in *minutes* seconds on a *personal computer* smartphone



# Why machine learning?

Machine learning can be:

- ▶ Powerful
  - ▶ With respect to computational efficiency
- ▶ Flexible
  - ▶ With respect to **data generating processes**
- ▶ Reduce the burden on the researcher
  - ▶ With respect to generating data...
  - ▶ and estimating models

But ML is not a panacea!

- ▶ Cannot solve problems of poor research design
- ▶ Increased flexibility *can* lead to poor interpretability

# Machine learning and social science

ML also introduces issues of its own:

## Twitter apologises for 'racist' image-cropping algorithm

**Users highlight examples of feature automatically focusing on white faces over black ones**

**Sarah Silverman sues OpenAI and Meta claiming AI training infringed copyright**

# Prediction and inference

The classic dichotomy when introducing ML:

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_{1i}$$

- ▶ **Inference:** estimating the size/direction of the (causal) relationship between variables ( $\hat{\beta}$  problems)
- ▶ **Prediction:** estimating the outcome, using the relationships between variables ( $\hat{y}$  problems)

These two facets are connected:

- ▶ Knowing the size/direction of (all) relationships  $\rightarrow$  predict the outcome
- ▶ But we rarely know the true model
- ▶ Sometimes we can get good at  $\hat{y}$  problems without knowing  $\hat{\beta}$

# Going down a philosophical rabbit hole

We might want to ask:

- ▶ Is there such a thing as a “true data generating process?” (see Grimmer et al 2021)
- ▶ Could a model yield perfect predictions without learning causal relationships?

More specifically for this course:

- ▶ Is the “=” in  $y = \beta x + \dots$  a convenience more than a scientific claim?
- ▶ What is the purpose of a model?

## $\hat{y}$ and $\hat{X}$ problems

We can think about where a prediction problem lies:

- ▶  $\hat{y}$  (dependent variable)
  - ▶ To predict the probability of revolution...
  - ▶ ... or the weather tomorrow
  - ▶ These are not necessarily *inferential* problems
- ▶  $\hat{X}$  (independent variables):
  - ▶ Dimensions of interest that may be important to our theory, but which are:
    - ▶ Latent (i.e. not directly observable)
    - ▶ Difficult to measure “by hand”
  - ▶ Use ML to make predictions over  $X$
  - ▶ These estimates are used *downstream* to test a theory

# Classification and prediction

Within both problems, we can distinguish two types:

- ▶ **Prediction** – estimating the value of a continuous variable (sometimes referred to as a “regression” problem) e.g.,
  - ▶ The ideology of a politician
  - ▶ The number of votes received by a candidate
- ▶ **Classification** – estimating which *class* of a category an observation belongs to, e.g.,
  - ▶ Party identity (Republican/Democrat/Libertarian/Independent)
  - ▶ The topic of a tweet (foreign/domestic/personal, positive/negative)
  - ▶ Recidivism

# Supervised vs unsupervised methods

## Supervised methods

- ▶ **Training data** are a set of labelled examples, where  $y$  is our target prediction
- ▶ We use these examples to "learn" the relationship between  $y$  and  $X$
- ▶ Then predict  $y$  for a *new* unlabeled dataset (i.e. where the target variable is not observed)

Learning the relationship:

$$\underbrace{\begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}}_{\mathbf{y}^{\text{TRAIN}}} \quad \underbrace{\begin{bmatrix} 3.3 & 1.1 & 0 \\ 2.7 & 0.8 & 0 \\ 1.8 & 0.1 & 1 \\ \vdots & \vdots & \vdots \\ 5 & 1.2 & 0 \end{bmatrix}}_{\mathbf{X}^{\text{TRAIN}}}$$

Predicting on new data

$$\mathbf{X}^{\text{TEST}} = \begin{bmatrix} 3.5 & 1.9 & 1 \\ 5.4 & 0.3 & 0 \\ 1.7 & 0.5 & 1 \end{bmatrix}$$



## Notation

# Algebra

Throughout the course, we will follow the notation set out in Burkov (2019):

- ▶  $\theta$  and  $x$  are scalar - i.e. a single number
  - ▶ E.g.,  $\theta = 3.141$ ,  $x = 1$  etc.
- ▶  $\boldsymbol{\theta}$  and  $\boldsymbol{x}$  are vectors - i.e. an ordered list of scalar values
  - ▶ E.g.,  $\boldsymbol{\theta} = [0.5, 3, 2]$
- ▶  $\boldsymbol{\Theta}$  and  $\mathbf{X}$  are matrices
  - ▶ E.g.,

$$\mathbf{X} = \begin{bmatrix} 1 & 5 \\ 24 & -3 \end{bmatrix}$$

- ▶  $\mathbf{x}^{(k)}$  is the  $k$ th column of  $\mathbf{X}$
- ▶  $\mathbf{x}_i$  is the  $i$ th row in matrix  $\mathbf{X}$
- ▶  $x_i^{(k)}$  is the  $k$ th element of the row vector  $\mathbf{x}_i$

# Probability notation

Let  $p$  denote a probability distribution *function* that returns the probability of an event or observation:

- ▶  $p(A) = 0.5$  means the probability of event  $A$  is 0.5
- ▶ All probabilities are bounded between 0 and 1
  - ▶  $0 \leq p(A) \leq 1$

Conditional probabilities means the probability of an event **given** the value of another variable

- ▶  $p(A|c) = 0.25$

# Probability rules

Probabilities have some nice features:

- ▶  $p(A \text{ and } B) = p(A)p(B|A)$
- ▶ If  $p(A) \perp p(B)$ , then  $p(B|A) = p(B)$

As a result:

- ▶ If  $p(A) \perp p(B)$ , then  $p(A \text{ and } B) = p(A)p(B)$
- ▶ These rules explain why the probability of two coin-flips is = 0.25:
  - ▶  $P(\text{Flip 1} = \text{Heads}) = 0.5$
  - ▶  $P(\text{Flip 2} = \text{Heads} | \text{Flip 1} = \text{Heads}) = P(\text{Flip 2} = \text{Heads})$
  - ▶  $P(\text{Flip 1} = \text{Heads} \text{ *and* Flip 2} = \text{Heads}) = 0.5 \times 0.5 = 0.25$

# Notation quiz

What are the following:

1.  $\mathbf{a}$
2.  $y_i$
3.  $\beta$
4.  $\beta$
5.  $\Theta$

If  $p(A) = 0.5$ ,  $p(B) = 0.1$ , and  $p(B|A) = 0.3$ :

6. Is  $p(A) \perp p(B)$ ?
7. What is  $p(A \text{ and } B)$ ?

## Maximum Likelihood Estimation (a gentle introduction)

# Searching the hypothesis space

A researcher wants to characterise an outcome as the linear combination of predictor variables:

$$\mathbf{y} = \beta_0 + \beta_1 \mathbf{x}^{(1)} + \dots + \beta_k \mathbf{x}^{(k)}$$

- ▶ We will set aside all inferential/theoretical concerns
- ▶ Focused on parsimonious description in a linear space

Since  $\mathbf{X}$  is fixed (it's the data we observe), we need to find the **best**  $\beta$ :

- ▶ Need some way of searching across all possible values (“**hypotheses**”) and finding the one that best *fits* our data
  - ▶ Statistical learning!

## Bayes Theorem (from a frequentist perspective)

$$\underbrace{P(A|B)}_{\text{Posterior}} = \frac{\overbrace{P(B|A)}^{\text{Likelihood}} \times \overbrace{P(A)}^{\text{Prior}}}{\underbrace{P(B)}_{\text{Evidence}}}$$

We can use Bayes formula to estimate the posterior probability of some parameter  $\theta$ :

$$p(\theta|\mathbf{X}) \propto p(\mathbf{X}|\theta) \times p(\theta),$$

where  $\mathbf{X}$  is the data.



## Likelihood function

Let's suppose that we have no prior knowledge over  $\theta$ , so we'll drop the prior and focus specifically on the likelihood:

$$\mathcal{L}(\theta) = p(\mathbf{X}|\theta)$$

*How would we calculate this?*

## Likelihood function

Let's suppose that we have no prior knowledge over  $\theta$ , so we'll drop the prior and focus specifically on the likelihood:

$$\mathcal{L}(\theta) = p(\mathbf{X}|\theta)$$

*How would we calculate this?*

$$\begin{aligned}\mathcal{L}(\theta) &= p(\mathbf{x}_1|\theta) \times p(\mathbf{x}_2|\theta) \times \dots \times p(\mathbf{x}_n|\theta) \\ &= \prod_{i=1 \dots N} p(\mathbf{x}_i|\theta)\end{aligned}$$

i.e. the product of the probability of each observations within  $\mathbf{X}$ , given  $\theta$ .

## Comparing likelihoods

Suppose we have two alternative values of  $\theta$ :  $\theta^{(1)}, \theta^{(2)}$ . We can calculate the likelihood *ratio* (LR) of these two possible parameter values:

$$LR = \frac{\mathcal{L}(\theta^{(1)})}{\mathcal{L}(\theta^{(2)})}$$

*If  $LR > 1$ , which set of parameters would we pick?*

## Comparing likelihoods

Suppose we have two alternative values of  $\theta$ :  $\theta^{(1)}, \theta^{(2)}$ . We can calculate the likelihood *ratio* (LR) of these two possible parameter values:

$$LR = \frac{\mathcal{L}(\theta^{(1)})}{\mathcal{L}(\theta^{(2)})}$$

*If  $LR > 1$ , which set of parameters would we pick?*

►  $\theta^{(1)}$

# Maximum likelihood estimation

We can generalise this for all possible values of  $\theta$ :

$$\arg \max_{\theta \in \Theta} \mathcal{L}(\theta) = \prod_{i=1 \dots N} p(\mathbf{x}_i | \theta)$$

i.e., from the set of all possible parameter values  $\Theta$ , find the parameter value that maximises the likelihood function.

Hence, **maximum** likelihood estimation.

- ▶ How we calculate  $p(\mathbf{x}_i | \theta)$  will depend on the functional form of the underlying distribution
- ▶ We'll explore this specifically with respect to logistic regression later on today

*Why is this maximum likelihood a useful concept?*

## Numeric overflow

Multiplying many small numbers means we soon lose the power to calculate them precisely

- ▶ R double-precision numbers range from  $2 \times 10^{-308}$  to  $2 \times 10^{308}$
- ▶ If 400 observations have  $p_{\theta} = 0.01$ ,  $\mathcal{L}(\theta)$  will be outside the computable range

What if we take the log?

- ▶ The log function is strictly increasing
- ▶  $\text{Log}(a \times b) = \text{Log}(a) + \text{Log}(b)$  so we can simply add the values

With the logged likelihood function we do not have the problem of numeric overflow!

# Negative log-likelihood

We can also calculate the *negative* log-likelihood:

- ▶ I.e. put a minus sign in front!
- ▶ We then *minimise* the negative log-likelihood
- ▶ We typically want to minimise rather than maximise because many of our procedures for optimisation are based on the former
- ▶ But, broadly, this is just semantics:
  - ▶ Minimising the negative log-likelihood is the same as maximising the log-likelihood

# Logistic regression

Logistic regression:

- ▶ Allows us to estimate  $\beta$  parameters when we have a binary outcome variable
- ▶ More broadly, it is a **binary classification** algorithm – what is the probability that  $y_i = 1$  given a vector of features  $\mathbf{x}_i$ ?

We can write the logistic regression function as,

$$f_{\theta,b}(\mathbf{X}) = \frac{1}{1 + e^{-(\theta\mathbf{X}+b)}}.$$

The goal is to find the *best* values of  $\theta$  and  $b$  that “explains” the data

- ▶ For simplicity, let's include  $b$  within  $\theta$  s.t.  $\theta = \{b, \theta_1, \dots, \theta_k\}$



# MLE of logistic regression I

For a given vector of scalar values  $\theta$ , we can ask what the likelihood of the data is given those values

- ▶ With the optimal parameter choice  $\theta^*$ :
  - ▶ When  $y_i = 1$ ,  $f_{\theta^*}(\mathbf{x}_i) = 1$
  - ▶ When  $y_i = 0$ ,  $f_{\theta^*}(\mathbf{x}_i) = 0$

*Why can't we just use the predicted probabilities as the likelihood?*

# MLE of logistic regression I

For a given vector of scalar values  $\theta$ , we can ask what the likelihood of the data is given those values

- ▶ With the optimal parameter choice  $\theta^*$ :
  - ▶ When  $y_i = 1$ ,  $f_{\theta^*}(\mathbf{x}_i) = 1$
  - ▶ When  $y_i = 0$ ,  $f_{\theta^*}(\mathbf{x}_i) = 0$

*Why can't we just use the predicted probabilities as the likelihood?*

- ▶ Predicted probabilities work well when  $y_i = 1$
- ▶ But we would erroneously down-weight the likelihood for all  $y_i = 0$

*How can we fix this?*

- ▶ When  $y_i = 0$ , let the likelihood equal  $1 - f_{\theta^*}(\mathbf{x}_i)$

## MLE of logistic regression II

We construct the likelihood for *any* datapoint using a mathematical “logic gate”:

$$\mathcal{L}_{\theta} = f_{\theta}(\mathbf{X})^y \times (1 - f_{\theta}(\mathbf{X}))^{(1-y)},$$

as when  $y_i = 0$ ,  $x^{y_i} = 1$  and  $x^{(1-y_i)} = x$ , and *vice versa*.

Simplifying, since  $f_{\theta}(\mathbf{x}_i) = \hat{y}_i$ :

$$\mathcal{L}_{\theta} = \hat{\mathbf{y}}^y (1 - \hat{\mathbf{y}})^{1-y}$$

# MLE optimization

We can then apply our “tricks” to make the computation easier:


$$-Log(\mathcal{L}_{\theta}) = -\sum_{i=1}^N Log(\mathcal{L}_{\theta}(\mathbf{x}_i)),$$

with the goal of minimising this quantity through choosing  $\theta$ .

How exactly do we minimize this function?

- ▶ Unlike OLS it is not possible to minimize this function analytically
- ▶ We therefore have to use computation to iterate through values of  $\theta$  to *approximate* the minima

## Minimising the negative log-likelihood in one dimension



# Gradient descent algorithm

To find the minimum of the negative log-likelihood we:

1. Choose a value for the starting parameter  $\theta$
2. Calculate the slope of the function at that point
3. Adjust our value of  $\theta$  in the *opposite* direction to the slope coefficient's sign
4. Recalculate the slope, and repeat 2-4

We can generalise this to  $\theta$ :

- ▶ Let  $Q(\theta)$  be the negative log likelihood function
- ▶ Calculate the **gradient vector** of the function in  $k$ -dimensions
- ▶ Adjust each parameter  $\theta_k \in \theta$  by the negative of the corresponding gradient element

$$\theta_k = \theta_k - \frac{\partial Q(\theta)}{\partial \theta_k}$$

# Logistic regression gradient


The partial derivative for any predictor  $\mathbf{x}^{(j)}$  for the *logistic* cost function is:

$$\frac{\partial Q^{\text{Logit}}}{\partial \theta_j} = (f_{\theta_j}(\mathbf{X}) - \mathbf{y}) \mathbf{x}^{(j)}$$

Hence the gradient of the function's curve for any vector of logistic parameters  $\boldsymbol{\theta}$  can be described as:

$$\nabla = \begin{bmatrix} \frac{\partial Q^{\text{Logit}}(\boldsymbol{\theta})}{\partial \theta_1} \\ \frac{\partial Q^{\text{Logit}}(\boldsymbol{\theta})}{\partial \theta_2} \\ \vdots \\ \frac{\partial Q^{\text{Logit}}(\boldsymbol{\theta})}{\partial \theta_k} \end{bmatrix}$$

# Progression of the descent algorithm





# Learning rate

As we iteratively adjust the value of our parameter:


- ▶ It's possible we keep jumping over the minima
- ▶ Or we get stuck in a rut and the estimator fails to find an even better parameter choice

So we can scale the impact of the current gradient on the new parameter choice:

- ▶ Let's call this scalar the **learning rate** ( $\lambda$ )
- ▶  $\theta_{\text{New}} = \theta - \lambda \nabla$

The choice of  $\lambda$  is down to the researcher:

- ▶ Overly-large values will not converge
- ▶ Overly small values may take too long, or risk converging on *local* minima



# Stochastic gradient descent

Gradient descent can be **expensive**:

- ▶ We have to evaluate all rows in our training data before making any updates to the parameters
- ▶ If we have lots of observations
  1. Each calculation takes a long time
  2. Takes many iterations to optimise
- ▶ Instead we can use **stochastic gradient descent** (SGD)
  - ▶ Inspect the loss of each observation (or a random subset) individually
  - ▶ Update the coefficients based on each observation

# Stochastic gradient descent

Under GD, for each iteration:

$$\theta_k \leftarrow \theta_k - \lambda \sum_{i=1}^N \left( f_{\theta_k}(\mathbf{x}_i) - y_i \right) \mathbf{x}_i^{(k)}$$

Under SGD, for each iteration:

$$\theta_k \leftarrow \theta_k - \lambda \left( f_{\theta_k}(\mathbf{x}_i) - y_i \right) \mathbf{x}_i^{(k)}, \text{ for } i \in \{1, \dots, N\}$$

- ▶ SGD typically converges a lot faster than GD
  - ▶ Every iteration we make  $N$  small changes to the parameter estimate
  - ▶ Computationally more efficient (we'll cover this more later in the week)
  - ▶ At the cost of some additional noise in the optimisation process

# Readings

Three suggested readings after today's class

- ▶ See course outline for more details
- ▶ Happy to answer questions on these tomorrow!

Coding workshop: writing our own logistic  
regression classifier